
Guarded Block Transfers for Polytope Distance: Linear Convergence with Explicit Constants and Gap-Certified Screening for SVM Training

Anonymous Author(s)

Affiliation

Address

email

Abstract

Hard-margin SVM training is the distance problem between two convex hulls, and the Triangle Algorithm (TA) solves it with a certified duality gap from supporting hyperplanes. We study the weight-transfer steps used by the enhanced TA and propose *guarded block transfers*: k Mitchell–Demyanov–Malozemov (MDM) transfers per scan are aggregated into a single exact line search, compared against the single best transfer at a cost independent of n , and replaced by an away step whenever the leading transfer is capacity-clipped. Analysing the method on the Minkowski-difference polytope, where the objective is 1-strongly convex and one pyramidal width governs the geometry, we prove global linear convergence with explicit constants and without the combinatorial swap-step bounds of pairwise Frank–Wolfe. Because the TA’s certified gap is the Frank–Wolfe gap over that polytope divided by the current distance, gap-safe screening applies, and we give an explicit activation theorem: screening removes every zero-weight non-support point after $O\left(\frac{k}{\rho} \log \frac{1}{\tau}\right)$ iterations, where ρ is the linear rate constant and τ the optimal-face score margin. Every inequality in the analysis has a numerical check. On five LIBSVM benchmarks the solver matches the baselines’ solutions to five digits wherever both converge. It wins or ties on sparse-support cells (gisette’s linear, kernel and L_2 margins; kernel L_2 margins with expensive columns) against LIBSVM and kernel sequential minimal optimisation (SMO), and it is one to two orders of magnitude slower on dense-support L_2 margins and nearly touching kernel hulls, where LIBLINEAR finishes in seconds. We report both sides.

1 Introduction

Training a linear hard-margin SVM is equivalent to finding the closest points of the two classes’ convex hulls: the optimal separator is the perpendicular bisector of the witness pair, and the margin is half the hull distance [8]. Kalantari’s Triangle Algorithm (TA) exploits this *distance duality* directly [1]. It comes in two forms: TA I decides whether the hulls intersect, and TA II computes their distance, maintaining an iterate pair (p, q) and a certified interval $[LB, UB]$ on the distance from parallel supporting hyperplanes, so that it can stop at a provable relative accuracy ε . The enhanced TA (ETA) of [2] adds joint segment updates, dot-product caching, an anti-zig-zag heuristic and prioritised search, and reports large speedups over SMO on synthetic data.

The workhorse step of such geometric solvers is the *weight transfer*: move mass from the worst active vertex to the best candidate vertex, the Mitchell–Demyanov–Malozemov (MDM) step [4], which is exactly the pairwise Frank–Wolfe (FW) step of [3] on the relevant polytope; MDM converges linearly [24], and FW methods have a long history in SVM training [20, 22, 23]. Each step

costs an $O(n)$ scan and produces one transfer; we ask how much of the scan can be reused across several transfers without weakening the guarantee.

Contributions.

1. **Guarded block transfers** (§3): pair the top- k receivers with the worst- k active donors, take one exact line search along the aggregated direction, and *guard* by a comparison with the single MDM step that costs $O(k^2)$ cached dot products, independent of n ; on capacity-clipped iterations, fall back to an away step. Cache updates for a whole block are one matrix–vector product.
2. **Linear convergence with explicit constants** (Theorem 1): on the Minkowski-difference polytope $Z = \text{conv}(V) - \text{conv}(W)$ the guarded block algorithm inherits the pyramidal-width rate of pairwise FW up to a constant, and the away fallback replaces pairwise FW’s combinatorial swap-step bound with linear drop-step counting, as blended pairwise conditional gradients do with a different guard [12].
3. **Screening activation with explicit constants** (Theorem 2): the TA certificate is the FW gap divided by the distance (Lemma 3, an observation that goes back to Gilbert and Wolfe), so gap-safe screening [16] applies, and we bound the iteration by which it has removed every zero-weight non-support point in terms of the rate constant and the optimal-face margin.
4. **A regime study on real data** (§5). The analysis predicts $O(n^2)$ -type cost when the optimal support is dense (pyramidal width $\sim 1/\sqrt{n}$, rate constant $\sim 1/n$). Two batteries on five LIBSVM datasets show the boundary from both sides: the solver wins or ties on sparse-support high-dimensional cells and loses by orders of magnitude on dense-support L_2 soft margins and on near-touching kernel hulls.

Code, proofs, numerical checks and experiment scripts are in the supplementary repository (anonymised for review).

2 Setting

Let $V = \{v_1, \dots, v_n\}$, $W = \{w_1, \dots, w_m\} \subset \mathbb{R}^d$. The polytope distance problem is

$$\min_{p \in \text{conv}(V), q \in \text{conv}(W)} \frac{1}{2} \|p - q\|^2. \quad (1)$$

Writing $z = p - q$ and $Z := \text{conv}(V) - \text{conv}(W) = \text{conv}\{v_i - w_j\}$, (1) is the minimum-norm-point problem $\min_{z \in Z} F(z)$, $F(z) = \frac{1}{2} \|z\|^2$, whose objective is 1-strongly convex and 1-smooth. Let z^* be the minimiser, $\delta^* = \|z^*\|$, $h(z) = F(z) - F(z^*)$, $M = \text{diam}(Z)$ and $\delta = \text{PW}(A) > 0$ the pyramidal width [3] of the atom set $A = \{v_i - w_j\}$ (a property of the nm atoms, including non-extreme ones, not of Z alone).

The algorithm keeps simplex weights α on V , β on W , with $p = \sum \alpha_i v_i$, $q = \sum \beta_j w_j$; the induced product representation $\lambda_{ij} = \alpha_i \beta_j$ of z over the atoms $v_i - w_j$ has support $\text{supp } \alpha \times \text{supp } \beta$. With $g = -\nabla F(z) = q - p$, the *side scores* are $s_i^V = \langle g, v_i \rangle$ and $s_j^W = \langle p - q, w_j \rangle$; the implementation maintains them incrementally as differences of cached arrays Vp, Vq, Wp, Wq .

Certificate. For any iterate, $\text{UB} = \|z\|$ and $\text{LB} = \min_{u \in Z} \langle z, u \rangle / \|z\| = (\min_i \langle z, v_i \rangle - \max_j \langle z, w_j \rangle) / \text{UB}$ satisfy $\text{LB} \leq \delta^* \leq \text{UB}$; the TA stops when $\text{UB} - \text{LB} \leq \varepsilon \text{UB}$. The L_2 soft margin with regularisation parameter C , the ν -SVM and the kernel variants reduce to (1) exactly (augmented coordinates that turn the Gram matrix K into $K + I/C$; reduced convex hulls; feature-space images), see Appendix B.

Gaps. $g^{\text{FW}} = \max_{u \in Z} \langle g, u \rangle - \langle g, z \rangle$, $g^A = \langle g, z \rangle - \min_{u \in \text{supp } \lambda} \langle g, u \rangle$, $g^{\text{PW}} = g^{\text{FW}} + g^A$. Because $\langle g, v_i - w_j \rangle = s_i^V + s_j^W$ is separable and the support is a product set, each gap splits into a V -side and a W -side term computable from the score arrays in $O(n + m)$ (Lemma 4), and a transfer $v_r - v_u$ on the V side is a pairwise direction of Z (Lemma 5).

Algorithm 1 One iteration of the guarded block-transfer step (side V ; W symmetric)

```

1: top- $k$  receivers by  $s^V$ ; worst- $k$  donors in  $\text{supp } \alpha$ ; pair best-with-worst:  $(r_j, u_j)$ ,  $j \leq k'$  ( $k' \leq k$ 
   pairs, limited by the support size); pair 1 is the MDM pair  $(r^*, u^*)$ 
2:  $d_j \leftarrow v_{r_j} - v_{u_j}$ ;  $\text{gap}_j \leftarrow s_{r_j}^V - s_{u_j}^V$ ; discard pairs with  $\text{gap}_j \leq 0$  and receivers that are donors;
    $\gamma_j \leftarrow \min(\text{gap}_j / \|d_j\|^2, c_j)$   $\triangleright c_j$ : donor capacity
3: if pair 1 uncapped ( $\gamma_1 = \text{gap}_1 / \|d_1\|^2$ ) then
4:    $\Delta_B \leftarrow$  gain of exact line search  $t \in [0, 1]$  along  $d = \sum_j \gamma_j d_j$ ;  $\Delta_1 \leftarrow$  gain of the MDM
   step
5:   take the step with the larger gain  $\triangleright$  guard,  $O(k'^2)$ 
6: else
7:   if the away step on  $u^*$  is boundary-clipped then take it  $\triangleright$  drop step:  $\alpha_{u^*} \rightarrow 0$ 
8:   else take the best of {block, MDM, away( $u^*$ ), toward( $r^*$ )}
9:   end if
10: end if

```

80 3 Guarded block transfers

81 Algorithm 1 amortises one $O(n)$ scan over up to k transfers. Only the scan is amortised: a new
 82 receiver still costs one Gram or kernel column, $O(nd)$, and a block touches up to $2k'$ columns, so
 83 Section 5 reports iterations, column evaluations and wall-clock separately. Every candidate gain is
 84 available from cached quantities at a cost independent of n : with $S = \langle g, d \rangle$ and $D = \|d\|^2$, the
 85 block gain is $tS - t^2 D/2$ at $t = \min(S/D, 1)$, where $S = \sum_j \gamma_j \text{gap}_j$ and $D = \gamma^\top \Gamma \gamma$ for the
 86 $k' \times k'$ Gram matrix Γ of the pair directions, $O(k'^2)$ dot products read off the cached columns; the
 87 MDM, away and toward gains are the standard one-dimensional expressions. Throughout, the *gain*
 88 of a step is the decrease $h(z_t) - h(z_{t+1})$ in the objective. The cache update for the whole block is
 89 a single matrix–vector product with the net per-index coefficients. Feasibility is immediate: donors
 90 lose at most $t\gamma_j \leq c_j$, receivers only gain, and the weight sum is conserved.

91 **Safe screening.** With $r = \sqrt{\text{UB}^2 - \max(\text{LB}, 0)^2}$ and $v_{\min} = \arg \min_i \langle z, v_i \rangle$, a zero-weight
 92 point with $\langle z, v_i \rangle - \langle z, v_{\min} \rangle > r \|v_i - v_{\min}\|$ is outside the optimal support and is removed (sym-
 93 metrically on W with the max). Because the optimum’s support survives, the reduced problem has
 94 the same optimum and all subsequent lower bounds remain valid.

95 4 Analysis

96 Full proofs are in Appendix A; we state the results and the ideas.

97 **Lemma 1** (Block gain). *Let $b_j = \gamma_j \|d_j\|$ over the k' retained pairs, all with $\text{gap}_j > 0$. Then*
 98 *$S \geq \sum_j b_j^2$, $D \leq k' S$, and the realised block gain satisfies $\Delta_B \geq S/(2k') \geq \Delta_{\text{MDM}}/(2k')$.*

99 The proof is the triangle inequality on $\|d\|$ and Cauchy–Schwarz on $(\sum b_j)^2$; capacity clipping only
 100 helps. With the guard, every case-(a) iteration’s gain is $\max(\Delta_B, \Delta_1) \geq \Delta_1$: the guarded step never
 101 makes less progress than pairwise FW.

102 **Lemma 2** (Good-step progress). *Work on the side with the larger pairwise gap. On any non-drop*
 103 *iteration, gain $\geq \min(g^{\text{PW}^2}/(32M^2), g^{\text{PW}}/8)$.*

104 Case (a) uses $g_{\text{side}}^{\text{PW}} \geq g^{\text{PW}}/2$ and the exact line search bound; case (b) uses $g_{\text{side}}^{\text{PW}} = g_{\text{side}}^{\text{FW}} + g_{\text{side}}^{\text{A}}$,
 105 so one of the toward step (always feasible with step ≤ 1) or the interior away step realises at least a
 106 quarter of g^{PW} .

107 **Theorem 1** (Global linear convergence). *Run the guarded block-transfer algorithm with away fall-*
 108 *back, one side per iteration (the side with the larger pairwise gap), any $k \geq 1$, from any ver-*
 109 *tex pair (initial support $s_0 \leq 2$), and write $h_t = h(z_t)$. Then every non-drop iteration satisfies*
 110 *$h_{t+1} \leq (1 - \rho)h_t$ with*

$$\rho = \frac{1}{16} (\delta/M)^2,$$

111 and among the first T iterations the number of drop steps is at most $kT_{\text{good}} + s_0$, where T_{good} is the
 112 number of non-drop iterations; hence $h_T \leq h_0(1 - \rho)^{(T-s_0)/(k+1)}$. The two-sided schedule of the
 113 implementation (larger-gap side first, second side skipped after a drop) satisfies the same statement
 114 with the same ρ and exponent $(T - s_0)/(2k + 1)$ (Lemma 7).

115 The rate follows from Lemma 2 and the geometric strong-convexity bound $g^{\text{PW}} \geq \delta\sqrt{2h}$ (Lemma 6,
 116 from the pyramidal-width definition of [3] and 1-strong convexity), together with $g^{\text{PW}} \geq g^{\text{FW}} \geq h$
 117 for the linear branch. The drop counting is the away-step argument: each drop removes one support
 118 index, non-drop steps add at most k , so drops are at most k times the good steps plus s_0 . There is
 119 no swap-step bound because case (b) never takes a capacity-clipped pairwise step: it takes an away
 120 step, which either contracts or drops. The implementation steps both sides per iteration with the
 121 schedule of Lemma 7 (larger-gap side first, second side skipped after a drop); the batteries of §5
 122 were run with the earlier unconditional V -then- W order, which differs only on drop iterations and
 123 agrees with it on the instances of Appendix C.

124 **Proposition 1** (Near-orthogonal speed-up). *If all k' pairs are uncapped, the block line search is*
 125 *interior, and $|\langle d_i, d_j \rangle| \leq \eta \|d_i\| \|d_j\|$ for $i \neq j$, then the block gain is at least $(\sum_j \Delta_j)/(1 + \eta(k' - 1))$,*
 126 *where Δ_j is pair j 's individual gain.*

127 As $\eta \rightarrow 0$ a single scan realises the sum of all k pairwise gains; in high dimension distinct support
 128 directions are nearly orthogonal, which is the mechanism behind the measured 4–30 $\times$ iteration
 129 reductions.

130 **Lemma 3** (The certified gap is the FW gap). *At any full scan, $\text{UB} - \text{LB} = g^{\text{FW}}(z)/\text{UB}$, and*
 131 *$g^{\text{FW}}(z) \leq 2h + R\sqrt{2h}$ with $R = \max_{u \in Z} \|u\|$.*

132 **Theorem 2** (Screening activation). *Assume $\delta^* > 0$, let F_V, F_W be the optimal faces (the vertices*
 133 *that can carry weight at the optimum), let $\tau_i = \langle z^*, v_i \rangle - \min_j \langle z^*, v_j \rangle$ be the score margin of v_i*
 134 *(and symmetrically on W), $\tau > 0$ the minimal nonzero margin, and L the larger class diameter.*
 135 *A zero-weight point with margin τ_i is screened as soon as $r < \tau_i/(3L)$; consequently, once $h_t <$
 136 $\tau^4 \delta^{*2}/(2592 R^4 L^4)$, which under Theorem 1 holds for all $t > s_0 + \frac{k+1}{\rho} \log(2592 R^4 L^4 h_0/(\tau^4 \delta^{*2}))$,
 137 every zero-weight point outside $F_V \cup F_W$ is removed permanently, and each later scan costs
 138 $O(|F_V| + |F_W| + |\text{supp } \lambda|)$ instead of $O(n + m)$.*

139 We expect the τ -dependence to be unavoidable, since a rule that is safe for every instance consistent
 140 with a certificate of gap $\sim \tau$ should not be able to distinguish a point of margin τ from a support
 141 point; we do not prove this. Each of Lemma 1, Proposition 1, Lemma 3, Lemma 7 and Theorem 2
 142 is verified numerically on random and adversarial ensembles at machine precision (Appendix E).

143 5 Experiments

144 **Setup.** All solvers are validated against an exact quadratic program (QP; SLSQP over the two
 145 simplices) on small instances: distances agree to 10^{-9} or better for every variant. Stopping rule
 146 $\varepsilon = 10^{-3}$ on the certified relative gap throughout unless stated; baselines use matched tolerances
 147 on the same objective (Appendix C). Synthetic instances follow the protocol of [2] (two Gaussian
 148 clouds, $n = 5,000$ per class). Real data are the LIBSVM sets a9a, w8a, ijcnn1, gisette and covtype
 149 (100k subsample), standardised, with official test sets where they exist, five seeds, and a 600 s budget
 150 per solver and cell. All real-data runs share one dedicated 32-vCPU Xeon 2.6 GHz machine (122 GB,
 151 OpenBLAS, one BLAS thread per solver, 26 cells concurrently), so times are comparable within a
 152 battery. Every ETA-versus-baseline pair solves the *same* objective; a differently regularised solver
 153 appears only as a labelled reference. “Our SMO” is our own SMO implementation on the same
 154 objective; K-ETA is the kernelised ETA.

155 **Synthetic (Appendix C).** Block transfers cut iterations 4–30 $\times$ relative to single MDM but wall-
 156 clock only 1.2–2.8 $\times$, since each block iteration touches more columns (Table 4); screening gives
 157 1.4–2.4 $\times$ at $\varepsilon = 10^{-5}$ and is neutral at 10^{-3} , as Theorem 2 predicts.

158 **Real data, sparse-support regime (Table 1).** TA I first certifies linear separability: gisette is
 159 separable and a9a, w8a, ijcnn1 and covtype are not (hulls intersect, certified in 0.2–42 s), a diagnostic
 160 none of the baselines produces. On gisette ($d=5,000$; 1,233 of 6,000 points in the support) ETA ties

Table 1: Regime battery on real data: hard margins and kernels, five seeds, 600 s per solver, mean $\pm$ 95% CI. Kernel cells (radial basis function kernel, RBF, $\gamma = 1/d$) use a class-balanced subsample of n_k training points. ETA and its baseline solve the same objective: LIN, linear hard margin (ETA with screening vs LIBSVM at $C=10^6$; our SMO in brackets); KHM, feature-space hard margin (K-ETA vs LIBSVM at $C=10^6$); KL2, kernel L_2 margin $K + I/C$, $C=1$ (K-ETA vs kernel SMO). conv = seeds in which ETA reached the certified gap (or certified intersection); oracle = $O(nd)$ column (ETA) or row (SMO) evaluations, in thousands. On datasets with fixed splits (all but the kernel subsamples) the CIs measure timing noise only.

cell	dataset	n_k / d	ETA (s)	baseline (s)	conv	dist. ETA / base	acc. ETA / base	oracle
LIN	gisette	6000 / 5000	56.6 $\pm$ 5.9	51.4 $\pm$ 13.1 [405 $\pm$ 62]	5/5	3.6638 / 3.6640	.981 / .981	1.5 [11.1]
KHM	gisette	6000 / 5000	86.2 $\pm$ 14.8	104.6 $\pm$ 20.2	5/5	.05215 / .05215	.979 / .979	2.6
KHM	ijcnn1	9853 / 22	421 $\pm$ 149	1.3 $\pm$ 0.3	4/5	.00430 / .00430	.946 / .946	2.5
KHM	a9a	10000 / 123	> 600	26.4 $\pm$ 2.9	0/5	$\leq$.00009 / .0033	.690 / .752	5.9
KHM	covtype	10000 / 54	> 600	144 $\pm$ 7 (maxiter)	0/5	$\leq$.00053 / .00025	.704 / .712	7.2
KHM	w8a	6479 / 300	101 $\pm$ 159	—	5/5	hulls intersect	—	1.9
KL2	gisette	6000 / 5000	131 $\pm$ 34	420 $\pm$ 49	5/5	.06882 / .06882	.971 / .971	3.6 / 15.2
KL2	w8a	6479 / 300	11.5 $\pm$ 2.6	41.7 $\pm$ 10.0	5/5	.05682 / .05682	.975 / .975	3.6 / 20.2
KL2	a9a	10000 / 123	52.8 $\pm$ 71.0	57.7 $\pm$ 23.2	5/5	.03019 / .03019	.795 / .795	7.6 / 43.6
KL2	covtype	10000 / 54	26.0 $\pm$ 4.2	17.6 $\pm$ 0.5	5/5	.02609 / .02609	.792 / .792	9.1 / 56.2
KL2	ijcnn1	9853 / 22	9.2 $\pm$ 2.5	6.1 $\pm$ 0.7	5/5	.04790 / .04790	.959 / .959	4.6 / 26.1

LIBSVM on the linear hard margin, both $7\times$ ahead of our SMO; K-ETA is $1.2\times$ faster than LIBSVM on the RBF hard margin and $3.2\times$ faster than kernel SMO on the kernel L_2 margin. Distances agree to five digits, support counts within three, test accuracies exactly. Across all five KL2 cells K-ETA returns the same distance and accuracy as kernel SMO with 4–6 $\times$ fewer kernel-column evaluations; in wall-clock it is 3–4 $\times$ faster where columns are expensive (gisette, w8a) and 1.5 $\times$ slower where they are cheap (ijcnn1 $d=22$, covtype $d=54$), because the $O(n)$ NumPy scan per iteration then dominates. The reference hinge-loss SVC ($C=1$, a different objective) takes 1–84 s with accuracies within one point (Appendix C).

The KHM column shows where the geometry turns against the method. On ijcnn1 both solvers return the identical solution (distance 0.00430, 819 support vectors, accuracy 0.946), but K-ETA needs $1.5\text{--}2.5 \times 10^5$ iterations and 420 s against LIBSVM’s 1.3 s: the feature-space hulls are 0.004 apart in a space of diameter $\approx \sqrt{2}$, so δ/M in Theorem 1 is of order 10^{-3} , and LIBSVM’s second-order working-set selection with heuristic shrinking is far better suited to that conditioning. On a9a and covtype the hulls are within 10^{-4} and 5×10^{-4} of touching: K-ETA times out at gap 3–10, while LIBSVM’s $C=10^6$ surrogate stops at its Karush–Kuhn–Tucker (KKT) tolerance at a point 6–30 $\times$ farther out than K-ETA’s upper bound (or hits its iteration cap on covtype), so neither solver certifies anything. K-ETA’s closer iterate has the lower test accuracy on a9a (0.690 vs 0.752) because the classifier is the bisector of $p - q$, whose direction is ill-determined when $\|p - q\| \rightarrow 0$; the hard-margin kernel model is the wrong one there anyway (KL2 accuracies 0.79 vs 0.70–0.75). On w8a the feature-space hulls intersect (duplicate points with opposite labels); K-ETA certifies this in 101 s and no hard margin exists.

Real data, L_2 soft margins (Table 2). The same boundary appears within one problem class. Gisette is separable, so its L_2 margin has a sparse support (1,233 of 6,000 points) and ETA converges to the certified gap in 30–45 s at every C , $7\times$ faster than our SMO on the same objective. LIBLINEAR run to the same primal accuracy (tolerance 10^{-6} ; at $C=1$: 0.14900 vs ETA’s 0.14893, inside ETA’s certified 8×10^{-4}) takes 8 min to 1.5 h, but that is the tolerance needed to match ETA’s primal, not a property of the problem; we did not time LIBLINEAR at its default tolerance (Appendix C) and therefore claim no wall-clock advantage over it on gisette. What ETA adds is the certificate: at $C=10$ LIBLINEAR’s progress-based stopping rule fires, even at tolerance 10^{-6} , at a primal 5% above ETA’s certified optimum (0.1566 vs 0.1490, test accuracy 0.977 vs 0.981), because on separable data the primal is nearly flat at large C ; the certificate does not fire early. On the four overlapping datasets both geometric solvers hit the 600 s cap on every cell while LIBLINEAR finishes in 0.3–25 s. At $C=0.1$ ETA’s iterate already has LIBLINEAR’s test accuracy on a9a, ijcnn1 and w8a (0.849, 0.918, 0.987) when the cap is reached with the certificate still at 0.01–0.6: the iterate converges long before the certificate closes. At $C \geq 1$ the augmented hulls ($K + I/C$) approach each other, the certified gap at the cap exceeds one and the iterate is not yet usable; on cov-

Table 2: L_2 soft-margin battery: five seeds (three for covtype at $C=0.1$), 600 s per solver for ETA and our SMO (our own implementation; LIBSVM does not solve this objective), LIBLINEAR at tolerance 10^{-6} (Appendix C) and unbounded; mean $\pm$ 95% CI, which on the fixed-split datasets (all but covtype) measures timing noise only. “cap” = timed out in every seed; ETA’s certified relative gap is reported at its stopping point (mean [min, max]); accuracies are those of the returned iterate. Two covtype $C=0.1$ seeds were still running when this draft was prepared.

dataset (n / d)	C	ETA (s)	SMO (s)	LIBLINEAR (s)	ETA gap at stop	acc. ETA / SMO / LIBLINEAR
gisette (6,000 / 5,000)	0.1	35.3 $\pm$ 6.8	243 $\pm$ 53	500 $\pm$ 41	8×10^{-4}	.981 / .981 / .981
	1	43.3 $\pm$ 5.3	295 $\pm$ 6	3916 $\pm$ 115	8×10^{-4}	.981 / .981 / .981
	10	33.2 $\pm$ 4.6	236 $\pm$ 36	5396 $\pm$ 197	10×10^{-4}	.981 / .981 / .977
a9a (32,561 / 123)	0.1	cap	cap	3.7 $\pm$ 0.8	0.15 [0.10, 0.23]	.849 / .845 / .849
	1	cap	cap	6.8 $\pm$ 1.7	1.26 [1.25, 1.28]	.810 / .780 / .849
	10	cap	cap	2.3 $\pm$ 0.5	1.83 [1.80, 1.86]	.635 / .763 / .849
ijcnn1 (49,990 / 22)	0.1	cap	478 $\pm$ 181	0.3 $\pm$ 0.1	0.010 [0.005, 0.016]	.918 / .918 / .918
	1	cap	cap	0.6 $\pm$ 0.2	0.68 [0.61, 0.73]	.901 / .917 / .918
	10	cap	cap	0.4 $\pm$ 0.1	1.34 [1.27, 1.40]	.656 / .852 / .918
w8a (49,749 / 300)	0.1	cap	cap	22.7 $\pm$ 4.3	0.57 [0.50, 0.66]	.987 / .985 / .987
	1	cap	cap	20.6 $\pm$ 6.6	1.17 [1.14, 1.21]	.923 / .979 / .987
	10	cap	cap	13.2 $\pm$ 1.1	2.02 [1.94, 2.14]	.720 / .971 / .987
covtype (100,000 / 54)	0.1	cap	cap	14.7 $\pm$ 22.8	1.04 [1.03, 1.05]	.542 / .651 / .755
	1	cap	cap	6.1 $\pm$ 3.6	1.42 [1.23, 1.73]	.508 / .541 / .756
	10	cap	cap	5.0 $\pm$ 1.8	2.01 [1.20, 2.37]	.517 / .514 / .756

type (100,000 points) it is not usable at any C . Uncapped, ETA does converge on the overlapping sets: on a9a ($C=0.1$, ten runs) in $2,908 \pm 138$ s and 95,201 iterations, with primal objective 687.07, identical to LIBLINEAR’s to five digits, and on w8a in 5,153 s. Figure 1 (appendix) shows these traces: the linear-rate signature of Theorem 1, but with 60–65% of the points in the support and near-uniform weights: the active atoms form a near-regular simplex on $s \approx 0.6n$ vertices, whose pyramidal width is $2/\sqrt{s}$ times its side [3], so $\rho \sim 1/s$, the iteration count grows linearly in n and, at $O(n)$ per iteration, the cost is $O(n^2)$ -type. Centroid initialisation and larger blocks did not help (the former increased the iteration count, the latter traded iterations for per-iteration cost).

6 Related work

Distance duality and the TA [1, 2]; MDM [4] and its linear rate [24]; Gilbert’s algorithm [5] and Wolfe’s method [25]; FW for polytope distance and SVMs [20, 21, 22, 23]; the nearest-point and reduced-hull view of soft margins [6, 8]; away steps [7] and the pyramidal-width analysis [3], whose pairwise step is our MDM pair. Swap-step-free pairwise variants [12, 13] and block-coordinate FW [9] are the closest algorithmic relatives; gap-safe screening [14, 16, 17] and active-set identification for FW [18, 19] the closest to Theorem 2, with [10, 11, 15] as its SVM ancestors. Appendix D details the differences.

7 Conclusion

Guarded block transfers give the Triangle Algorithm a linear rate with explicit constants and no combinatorial swap-step bounds, together with a screening rule that is both safe and provably effective. The cost is a factor $(k + 1)$ in the exponent, recovered in the near-orthogonal regime where a block realises close to the sum of its k pairwise gains. The experiments locate the boundary the analysis predicts: sparse support and well-separated hulls, where the solver matches or beats LIBSVM and kernel SMO and supplies a certificate, against dense support or nearly touching hulls, where δ/M is small and primal coordinate descent or second-order SMO is faster by one to two orders of magnitude. Which side a run is on can be read from its certified gap and support size. Limitations: screening pays only at tolerances tight enough for the certificate to close; the implementation is NumPy-level (a compiled inner loop would move the practical boundary, not the $O(n^2)$ asymptotics) and does not support sparse inputs; the kernel experiments use 10k-point subsamples; the reduced-hull (ν -SVM) variant is validated but not analysed; and the constants in Theorem 1 are loose.

References

- [1] B. Kalantari. A characterization theorem and an algorithm for a convex hull problem. *Annals of Operations Research*, 226:301–349, 2015.
- [2] M. Gupta and B. Kalantari. An enhanced triangle algorithm for large-scale support vector machine optimization. *JIDMIS*, 3(9s), 2026.
- [3] S. Lacoste-Julien and M. Jaggi. On the global linear convergence of Frank–Wolfe optimization variants. *NeurIPS*, 2015.
- [4] B. F. Mitchell, V. F. Demyanov, and V. N. Malozemov. Finding the point of a polyhedron closest to the origin. *SIAM J. Control*, 12(1):19–26, 1974.
- [5] E. G. Gilbert. An iterative procedure for computing the minimum of a quadratic form on a convex set. *SIAM J. Control*, 4(1):61–80, 1966.
- [6] S. S. Keerthi, S. K. Shevade, C. Bhattacharyya, and K. R. K. Murthy. A fast iterative nearest point algorithm for support vector machine classifier design. *IEEE Trans. Neural Networks*, 11(1):124–136, 2000.
- [7] J. Guélat and P. Marcotte. Some comments on Wolfe’s ‘away step’. *Mathematical Programming*, 35:110–119, 1986.
- [8] K. P. Bennett and E. J. Brendensteiner. Duality and geometry in SVM classifiers. *ICML*, 2000.
- [9] S. Lacoste-Julien, M. Jaggi, M. Schmidt, and P. Pletscher. Block-coordinate Frank–Wolfe optimization for structural SVMs. *ICML*, 2013.
- [10] L. El Ghaoui, V. Viallon, and T. Rabbani. Safe feature elimination in sparse supervised learning. *Pacific J. Optimization*, 8(4):667–698, 2012.
- [11] K. Ogawa, Y. Suzuki, and I. Takeuchi. Safe screening of non-support vectors in pathwise SVM computation. *ICML*, 2013.
- [12] K. Tsuji, K. Tanaka, and S. Pokutta. Pairwise conditional gradients without swap steps and sparser kernel herding. *ICML*, 2022.
- [13] F. Rinaldi and D. Zeffiro. Avoiding bad steps in Frank–Wolfe variants. *Computational Optimization and Applications*, 84:225–264, 2023.
- [14] O. Fercoq, A. Gramfort, and J. Salmon. Mind the duality gap: safer rules for the Lasso. *ICML*, 2015.
- [15] J. Wang, P. Wonka, and J. Ye. Scaling SVM and least absolute deviations via exact data reduction. *ICML*, 2014.
- [16] E. Ndiaye, O. Fercoq, A. Gramfort, and J. Salmon. Gap safe screening rules for sparsity enforcing penalties. *JMLR*, 18(128):1–33, 2017.
- [17] A. Shibagaki, M. Karasuyama, K. Hatano, and I. Takeuchi. Simultaneous safe screening of features and samples in doubly sparse modeling. *ICML*, 2016.
- [18] I. M. Bomze, F. Rinaldi, and D. Zeffiro. Active set complexity of the away-step Frank–Wolfe algorithm. *SIAM J. Optim.*, 30(3):2470–2500, 2020.
- [19] D. Garber. Revisiting Frank–Wolfe for polytopes: strict complementarity and sparsity. *NeurIPS*, 2020.
- [20] B. Gärtner and M. Jaggi. Coresets for polytope distance. *SoCG*, 2009.
- [21] K. L. Clarkson. Coresets, sparse greedy approximation, and the Frank–Wolfe algorithm. *ACM Trans. Algorithms*, 6(4):63, 2010.
- [22] I. W. Tsang, J. T. Kwok, and P.-M. Cheung. Core vector machines: fast SVM training on very large data sets. *JMLR*, 6:363–392, 2005.
- [23] R. Nanculef, E. Frandi, C. Sartori, and H. Allende. A novel Frank–Wolfe algorithm. Analysis and applications to large-scale SVM training. *Information Sciences*, 285:66–99, 2014.
- [24] J. López and J. R. Dorronsoro. Linear convergence rate for the MDM algorithm for the nearest point problem. *Pattern Recognition*, 48(4):1510–1522, 2015.
- [25] P. Wolfe. Finding the nearest point in a polytope. *Mathematical Programming*, 11:128–149, 1976.

277 **AI assistance.** AI tools were used in the preparation of the code, experiments and manuscript; the
 278 authors verified all results and are responsible for the content.

279 A Proofs

280 **Lemma 4** (Separability). *With the product representation, $g^{\text{PW}} = g_V^{\text{PW}} + g_W^{\text{PW}}$ where $g_V^{\text{PW}} =$
 281 $\max_i s_i^V - \min_{i \in \text{supp } \alpha} s_i^V \geq 0$ and g_W^{PW} likewise; the same holds for g^{FW} and g^{A} .*

282 *Proof.* $\langle g, v_i - w_j \rangle = s_i^V + s_j^W$ is separable and the support is a product set, so the max over
 283 atoms and the min over the support separate; each side term is nonnegative since the support max
 284 dominates the support min and the global max dominates the support max. $\square$

285 **Lemma 5** (Side transfers are Z -pairwise directions). *For $\alpha_u > 0$ and any $b \in \text{supp } \beta$, $v_r - v_u =$
 286 $(v_r - w_b) - (v_u - w_b)$ is a difference of an atom of Z and an atom in the support, and the transfer
 287 $\gamma \leq \alpha_u$ realises $z \leftarrow z + \gamma(v_r - v_u)$ inside Z .*

288 *Proof of Lemma 1.* Since $\gamma_j \leq \text{gap}_j / \|d_j\|^2$ we have $\gamma_j \text{gap}_j \geq \gamma_j^2 \|d_j\|^2 = b_j^2$, so $S =$
 289 $\sum_j \gamma_j \text{gap}_j \geq \sum_j b_j^2$. By the triangle inequality $\|d\| \leq \sum_j \gamma_j \|d_j\| = \sum_j b_j$ and by Cauchy-
 290 Schwarz $(\sum_j b_j)^2 \leq k' \sum_j b_j^2 \leq k' S$, hence $D \leq k' S$. Exact line search on the 1-smooth F
 291 gives $\Delta_B = S^2 / (2D) \geq S / (2k')$ in the interior case and $\Delta_B = S - D/2 \geq S/2$ at the bound-
 292 ary $t = 1$ (where $D \leq S$). Finally $S \geq \gamma_1 \text{gap}_1 \geq \Delta_{\text{MDM}}$ since all terms are nonnegative and
 293 $\Delta_{\text{MDM}} = \gamma_1 \text{gap}_1 - \gamma_1^2 \|d_1\|^2 / 2$. $\square$

294 *Proof of Lemma 2.* Let $g_S = \max(g_V^{\text{PW}}, g_W^{\text{PW}}) \geq g^{\text{PW}} / 2$. Every direction below is a difference of
 295 two points of one class hull, so its norm is at most $\text{diam}(\text{conv } V) \leq M$. For a feasible direction
 296 with directional derivative $\sigma > 0$, squared norm D and maximal step $\bar{t}$, exact line search on a 1-
 297 smooth function gives gain $\geq \min(\sigma^2 / (2D), \bar{t}\sigma/2)$ (interior optimum, or the clipped step evaluated
 298 with $\bar{t}D \leq \sigma$). Case (a): the guard's gain is at least the MDM gain, whose derivative is g_S and
 299 whose interior optimum is attained, so gain $\geq g_S^2 / (2M^2) \geq g^{\text{PW}2} / (8M^2)$. Case (b), non-drop:
 300 $g_S = g_{\text{side}}^{\text{FW}} + g_{\text{side}}^{\text{A}}$, so one term is $\geq g_S / 2 \geq g^{\text{PW}} / 4$. If it is the FW term, the toward step (maximal
 301 step 1) gives gain $\geq \min(g_{\text{side}}^{\text{FW}2} / (2M^2), g_{\text{side}}^{\text{FW}} / 2) \geq \min(g^{\text{PW}2} / (32M^2), g^{\text{PW}} / 8)$; if it is the away
 302 term, the away line search is interior (non-drop) and gain $\geq g_{\text{side}}^{\text{A}2} / (2M^2) \geq g^{\text{PW}2} / (32M^2)$. The
 303 algorithm takes the best available candidate. $\square$

304 **Lemma 6** (Geometric strong convexity; [3, Thm. 6] with $\mu = 1$). *For every $z \in Z$ with a represen-*
 305 *tation whose support contains an away atom, $g^{\text{PW}}(z) \geq \delta \sqrt{2h(z)}$ and $g^{\text{PW}}(z) \geq g^{\text{FW}}(z) \geq h(z)$.*

306 *Proof.* This is Theorem 6 of [3] for a 1-strongly convex objective; we include the short argument.
 307 By the definition of the pyramidal width [3, Def. 3], for $e = (z^* - z) / \|z^* - z\|$ the FW atom s and
 308 the away atom v satisfy $\langle g, s - v \rangle \geq \delta \langle g, e \rangle$. By 1-strong convexity, $\langle g, z^* - z \rangle \geq h + \frac{1}{2} \|z^* - z\|^2$,
 309 so with $\varrho = \|z^* - z\|$, $g^{\text{PW}} \geq \delta(h/\varrho + \varrho/2) \geq \delta\sqrt{2h}$ by AM-GM. Finally $g^{\text{PW}} = g^{\text{FW}} + g^{\text{A}} \geq$
 310 $g^{\text{FW}} = \max_u \langle g, u - z \rangle \geq \langle g, z^* - z \rangle \geq h$ by convexity. $\square$

311 *Proof of Theorem 1.* By Lemmas 2 and 6 a non-drop iteration gains at least
 312 $\min(2\delta^2 h / (32M^2), h/8) = (\delta/M)^2 h / 16$, using $\delta \leq M$; hence $h_{t+1} \leq (1 - \rho)h_t$. Drop
 313 counting: non-drop iterations add at most k support indices, drop iterations add none and remove at
 314 least one, and the support size stays ≥ 2 from s_0 ; so the number of removals, hence of drops, is at
 315 most $kT_{\text{good}} + s_0$, giving $T_{\text{good}} \geq (T - s_0) / (k + 1)$. All steps, drops included, are descent steps,
 316 so h is monotone. $\square$

317 **Lemma 7** (Two-sided schedule). *Let every iteration compute both side gaps, take the guarded step*
 318 *of Algorithm 1 on the side with the larger pairwise gap first, and take the other side's guarded step*
 319 *only if the first step was not a drop; an iteration counts as a drop iff its first step is a drop. Then*
 320 *every non-drop iteration gains at least $\min(g^{\text{PW}2} / (32M^2), g^{\text{PW}} / 8)$ and $\# \text{drops} \leq 2kT_{\text{good}} + s_0$;*
 321 *hence Theorem 1 holds for this schedule with the same ρ and exponent $(T - s_0) / (2k + 1)$.*

322 *Proof.* On a non-drop iteration the first step is exactly the step analysed in Lemma 2, so it gains at
 323 least the stated amount; the second step is an exact line search along a feasible descent direction (or
 324 is skipped), so it cannot increase h . Drop counting: a drop iteration takes a single away step that
 325 removes one support index and adds none, because the second side is skipped; a non-drop iteration
 326 takes two guarded steps and adds at most k indices each. Removals, hence drops, therefore number
 327 at most $2kT_{\text{good}} + s_0$. $\square$

328 *Remark 1.* The skip is what makes the count work. If the second side always steps, an iteration
 329 whose analysed side drops can add up to k indices on the other side, and the number of such iterations
 330 is then bounded only by $\frac{k}{k+1}T$, which gives no rate for $k \geq 1$.

331 *Proof of Proposition 1.* $D = \sum_{i,j} \gamma_i \gamma_j \langle d_i, d_j \rangle \leq \sum_j b_j^2 + \eta \sum_{i \neq j} b_i b_j \leq (1 + \eta(k' - 1)) \sum_j b_j^2 \leq$
 332 $(1 + \eta(k' - 1))S$ by Cauchy–Schwarz on the cross terms; then $\Delta_B = S^2/(2D) \geq S/(2(1 + \eta(k' -$
 333 $1)))$, and for uncapped pairs $S = \sum_j 2\Delta_j$. $\square$

334 *Proof of Lemma 3.* $\text{LB} = \min_{u \in Z} \langle z, u \rangle / \|z\|$ by definition of the supporting hyperplane bound and
 335 $\text{UB} = \|z\|$; subtracting, $\text{UB} - \text{LB} = (\langle z, z \rangle - \min_u \langle z, u \rangle) / \|z\| = g^{\text{FW}} / \text{UB}$. At the optimum
 336 $0 = \langle z^*, z^* \rangle - \min_u \langle z^*, u \rangle$; subtracting from $g^{\text{FW}}(z)$, $g^{\text{FW}}(z) = (\|z\|^2 - \|z^*\|^2) + (\min_u \langle z^*, u \rangle -$
 337 $\min_u \langle z, u \rangle) \leq 2h + \max_u \langle z^* - z, u \rangle \leq 2h + R\|z - z^*\| \leq 2h + R\sqrt{2h}$ by 1-strong convexity. $\square$

338 *Proof of Theorem 2.* Strong convexity gives $\|z - z^*\| \leq r$ for the screening radius $r =$
 339 $\sqrt{\text{UB}^2 - \max(\text{LB}, 0)^2}$. Let $m^* = \min_j \langle z^*, v_j \rangle$ and pick $v' \in F_V$; since $v_{\min}$ minimises $\langle z, \cdot \rangle$,
 340 $\langle z^*, v_{\min} \rangle - m^* = \langle z, v_{\min} - v' \rangle + \langle z^* - z, v_{\min} - v' \rangle \leq rL$. Hence $\langle z, v_i - v_{\min} \rangle =$
 341 $\langle z^*, v_i - v_{\min} \rangle + \langle z - z^*, v_i - v_{\min} \rangle \geq \tau_i - rL - r\|v_i - v_{\min}\|$, and the rule fires when this
 342 exceeds $r\|v_i - v_{\min}\|$, for which $\tau_i > 3rL$ suffices since $\|v_i - v_{\min}\| \leq L$. For the activation
 343 time, $r^2 = (\text{UB} - \text{LB})(\text{UB} + \text{LB}) \leq 2R(\text{UB} - \text{LB}) = 2Rg^{\text{FW}}/\text{UB} \leq (2R/\delta^*)(2h + R\sqrt{2h}) \leq$
 344 $(4R^2/\delta^*)\sqrt{2h}$ for $h \leq R^2/2$ (Lemma 3); requiring $r^2 \leq \tau^2/(9L^2)$ yields the stated condition on
 345 h_t , and Theorem 1 with $\log(1/(1 - \rho)) \geq \rho$ converts it into the iteration bound. Permanence follows
 346 from the safety of the rule. $\square$

347 B Exact reductions

348 **L_2 soft margin.** The problem $\min \frac{1}{2}\|w\|^2 + \frac{C}{2} \sum_i \xi_i^2$ is the hard-margin problem on the augmented
 349 points $\tilde{x}_i = [x_i; e_i/\sqrt{C}]$, whose Gram matrix is $K + I/C$ (K the Gram matrix of the original points)
 350 and whose augmented hulls never intersect; let δ_C denote their distance. The augmentation is never
 351 materialised: the extra coordinates of an iterate are exactly its convex weights, so every augmented
 352 inner product is a cached original-space product plus a sparse diagonal correction, $\langle \tilde{p}, \tilde{v}_i \rangle = \langle p, v_i \rangle +$
 353 α_i/C and $\|\tilde{p}\|^2 = \|p\|^2 + \|\alpha\|^2/C$. The separator is recovered as $w = 2(p - q)/\delta_C^2$, $b = (\|\tilde{q}\|^2 -$
 354 $\|\tilde{p}\|^2)/\delta_C^2$, and the primal value is $P = 2/\delta_C^2$. Validation: the solver agrees with an exact QP on the
 355 explicitly augmented problem to 10^{-10} across random instances and values of C ; the regularisation
 356 path δ_C is monotone in C ; on a9a ($C=0.1$) its certified primal 687.07 equals LIBLINEAR’s to five
 357 digits.

358 **ν -SVM.** The L_1 (hinge) soft margin is the nearest-point problem between *reduced* convex hulls
 359 $R(V, \mu) = \{\sum a_i v_i : \sum a_i = 1, 0 \leq a_i \leq \mu\}$, $\mu = 2/(\nu\ell)$ with $\ell = n + m$ the number of
 360 training points [8]. Two changes: the extreme-point oracle becomes a capped top- k blend (weight μ
 361 on the $\lfloor 1/\mu \rfloor$ best points in the search direction, via a partial sort, still $O(n)$), and the primary step is
 362 the capped MDM transfer with blended toward steps as fallback; caching and the reduced support-
 363 function lower bound carry over unchanged. Validation: agreement with the exact box-constrained
 364 QP to 10^{-9} – 10^{-14} and with LIBSVM’s NuSVC to $\sim 10^{-10}$ through the mapping $\delta^* = \|w_{\text{svc}}\|/\varsigma$, ς
 365 the per-class sum of the rescaled dual coefficients. The reduced hull is a polytope, so (1) still applies,
 366 but its vertices are blends and the capped transfer is not the pairwise step analysed in §4; the analysis
 367 of this variant is left open.

Kernels. The algorithm touches data only through inner products, so kernelisation changes plumbing, not algorithm: Gram columns become kernel columns $\kappa(V, x_i)$, squared norms become kernel diagonals, iterates exist purely as convex weights (the exact cache refresh is reconstructed from the support’s cached kernel columns), and the classifier is the kernel expansion of the witness-pair bisector. A ridge $K + I/C$ gives the kernel L_2 margin by the same diagonal trick. Validation: the linear kernel reproduces the Euclidean solver to 10^{-12} ; RBF distances match the exact kernel QP to 10^{-9} and LIBSVM’s hard-margin solution, via $\|w_{\mathcal{H}}\|^2 = (\alpha \circ y)^\top K(\alpha \circ y)$ for LIBSVM’s dual coefficients α and labels y , to 10^{-5} ; kernel SMO matches the QP to 10^{-8} .

C Experimental details

Solver settings and tolerance equivalence. ETA: block transfers with $k = 16$, guarded, away fallback, pairwise anti-zig-zag, stop at certified relative gap 10^{-3} , iteration cap 2×10^6 , 600 s cap. The batteries were run with the unconditional V -then- W order; the schedule of Lemma 7 (larger gap first, skip after a drop) differs from it only on drop iterations, and Table 3 shows the two agree on the synthetic instances of this appendix; screening (every 25 iterations) only in the LIN cell, where it never activated at 10^{-3} , as Theorem 2 predicts. Our SMO: maximal-violating-pair selection, KKT gap 10^{-3} , same caps; hard margin as $C = 10^{12}$; kernel SMO on $K + I/C$. LIBLINEAR (scikit-learn 1.8 LinearSVC): squared hinge with parameter $C/2$, so that its objective equals $\frac{1}{2}\|w\|^2 + \frac{C}{2} \sum \xi_i^2$; tolerance 10^{-6} , chosen so that its final primal matches ETA’s certified value on every cell where ETA converges (a9a 687.07 vs 687.07; gisette 0.14900 vs 0.14893, inside ETA’s certified 8×10^{-4} ; ijcnn1 550.6 vs 550.5 at ETA gap 0.01); this tolerance is also why LIBLINEAR needs 18,000 iterations and over an hour on gisette at $C=1$, and 35,000 iterations and 1.5 h at $C=10$, where it nevertheless stops 5% above the certified optimum (0.1566 vs 0.1490). The gisette LIBLINEAR times in Table 2 are therefore the cost of matching ETA’s primal, not LIBLINEAR’s cost at its default tolerance (10^{-4}), which we did not measure; since its dual coordinate descent is monotone, the primal at the default tolerance can only be at or above the 10^{-6} value. The paper claims no wall-clock advantage over LIBLINEAR on gisette. Intercept scaling 100; iteration cap 2×10^5 . LIBSVM (scikit-learn SVC): tolerance 10^{-3} , 1000 MB cache, iteration cap 2×10^6 (reported as *maxiter* when hit); hard margins as $C = 10^6$; RBF with $\gamma = 1/d$. Oracle counts: ETA counts Gram or kernel columns evaluated against a full class ($O(nd)$ or $O(n_k d)$ each); SMO counts kernel rows, the same unit.

Data and seeds. a9a, w8a, ijcnn1 and gisette use the LIBSVM training and test files; covtype (binary, scaled) is subsampled to 100,000 points and split 75/25 by seed. Features are standardised on the training set. Kernel cells use a class-balanced subsample of at most 5,000 points per class (w8a’s minority class has 1,479 points, hence $n_k = 6,479$; gisette uses all 6,000). Seeds vary the covtype split, the kernel subsamples and the timing environment; on datasets with fixed splits the deterministic solvers return identical solutions across seeds, which is why several accuracy CIs are zero. Confidence intervals are t -based 95% intervals over seeds.

ETA on real data (L_2 soft margin, $C = 0.1$): certified gap to the 10^{-3} stopping rule




Figure 1: Certified relative gap of ETA on a9a and w8a (L_2 , $C = 0.1$, uncapped), versus iteration and wall-clock: linear-rate convergence in the dense-support regime, at $O(n^2)$ -type cost.

Hardware and concurrency. One dedicated 32-vCPU Intel Xeon 2.6 GHz Lightning Studio, 122 GB RAM, Linux, Python 3.12, NumPy 2.4 with OpenBLAS 0.3.31, OMP_NUM_THREADS=1; the

406 L_2 battery ran 14 cells and the regime battery 12 cells concurrently (the gisette, covtype and w8a
407 $C=10$ cells of Table 2 were completed in later passes with 11–14 workers and no other load). The
408 uncapped a9a/w8a runs in Figure 1 were made earlier on the same machine under a different load
409 (LIBLINEAR took 12 and 48 s there against 4 and 23 s in the capped battery), so ratios across the
410 two runs carry a factor of roughly two of uncertainty; ratios within a battery do not.

411 **Reference solver in the KL2 cells.** The hinge-loss SVC(rbf, $C=1$), a different objective, took
412 8.6 ± 2.9 (a9a), 3.4 ± 0.3 (covtype), 83.5 ± 13.1 (gisette), 1.1 ± 0.0 (ijcnn1) and 4.9 ± 1.2 s (w8a),
413 with test accuracies 0.789, 0.783, 0.977, 0.950 and 0.975, within one point of the L_2 solutions and
414 on either side of them.

415 **Step ablation, blocks and shrinking (synthetic).** Iteration counts at $\varepsilon = 10^{-6}$ (single instances):
416 $d=30$, $n=200$: toward 158,941, midpoint 148,996, away 106, pairwise 181; $d=100$, $n=1,000$: to-
417 ward and midpoint fail within 300,000, away 15,886, pairwise 3,921. At the working tolerance 10^{-3}
418 the pairwise/away gain over the midpoint heuristic is only $\sim 11\%$, which is why the heuristic suf-
419 ficed in [2]. Block transfers cut iterations 4–30 $\times$ relative to single MDM (e.g. 43,486 $\rightarrow$ 1,551 on a
420 dense-support L_2 instance at $d=1,000$) and are the only configuration that reaches the certified toler-
421 ance on MNIST odd-vs-even. Figure 2 shows the screening speed-ups and Figure 3 the consolidated
422 comparison: hard margins 0.24 vs LIBSVM’s 0.39 s at $d=1,000$ and 8.6 vs 20.5 s at $d=10,000$;
423 0.21 vs 0.42 s at $\varepsilon=10^{-5}$ (where the paper-style configuration fails to converge); RBF hard margin
424 0.76 vs 1.28 s; ν -SVM 18.5 vs 21.9 s; synthetic L_2 ($d=1,000$, $C=1$) 8.9 vs LIBLINEAR’s 13.8 s;
425 MNIST odd-vs-even 82.8 vs LIBLINEAR’s 55.0 s. Distances agree within tolerance on every row.

Table 3: Two-sided schedules on the synthetic instances of this appendix ($n=5,000$ per class, $k=16$): the unconditional V -then- W order used for the batteries versus the schedule of Lemma 7 (larger-gap side first, second side skipped after a drop). Drop steps and skipped second-side steps are counted over the run.

instance	schedule	time (s)	iterations	drops	skips	distance
hard margin, $d=1000$, $\varepsilon=10^{-5}$	unconditional	0.1	16	0	0	55.339365
hard margin, $d=1000$, $\varepsilon=10^{-5}$	Lemma 7	0.2	16	0	0	55.339365
L_2 overlap, $d=100$, $C=1$	unconditional	7.9	11,216	638	0	0.078400
L_2 overlap, $d=100$, $C=1$	Lemma 7	8.1	10,701	599	316	0.078400
L_2 overlap, $d=1000$, $C=1$	unconditional	4.3	2,951	172	0	0.578143
L_2 overlap, $d=1000$, $C=1$	Lemma 7	4.3	3,001	180	121	0.578143

Table 4: Block-size ablation on the same instances (current code): single MDM ($k=1$) versus blocks of k transfers. Iterations fall with k ; wall-clock does not fall in proportion, because each block iteration touches more columns.

instance	k	time (s)	iterations	columns	drops
hard margin, $d=1000$, $\varepsilon=10^{-5}$	1	0.3	296	104	0
hard margin, $d=1000$, $\varepsilon=10^{-5}$	4	0.2	36	103	0
hard margin, $d=1000$, $\varepsilon=10^{-5}$	16	0.2	16	103	0
hard margin, $d=1000$, $\varepsilon=10^{-5}$	32	0.2	16	109	0
L_2 overlap, $d=100$, $C=1$	1	10.7	75,021	1,343	0
L_2 overlap, $d=100$, $C=1$	4	11.2	24,606	1,632	416
L_2 overlap, $d=100$, $C=1$	16	9.0	10,701	2,404	599
L_2 overlap, $d=100$, $C=1$	32	10.4	7,606	3,076	305
L_2 overlap, $d=1000$, $C=1$	1	12.2	43,521	1,251	0
L_2 overlap, $d=1000$, $C=1$	4	6.7	10,071	1,330	177
L_2 overlap, $d=1000$, $C=1$	16	4.3	3,001	1,348	180
L_2 overlap, $d=1000$, $C=1$	32	4.6	1,741	1,396	170

426 D Extended related work

427 **Swap steps.** Blended pairwise conditional gradients [12] also remove pairwise FW’s swap-step
428 bound with a guard, there between the pairwise step and a FW step, and [13] avoid bad steps with
429 a short-step chain. Our guard compares an aggregated block of k transfers with the single MDM




Figure 2: Safe screening: wall-clock speed-up vs dimension at two tolerances on the separated synthetic protocol; labels are surviving points of 10,000.




Figure 3: Time to certified solution on seven problem classes: optimised ETA vs the paper-style configuration vs the best standard solver per class.

430 step and falls back to an away step, so the drop-step count of away-step FW applies unchanged.
 431 Block-coordinate FW [9] blocks over *coordinates* of a product domain rather than aggregating pairs
 432 on one polytope.

433 **Screening and identification.** Gap-safe rules [14, 16, 17] screen with a radius from the duality gap
 434 and identify the support in finite time; the active-set complexity of away-step FW [18] and the strict-
 435 complementarity bounds of [19] give the same identification from the algorithm’s side. Theorem 2
 436 is the polytope-distance instance, with constants tied to the TA certificate; the SVM rules of [10, 11]
 437 and the exact data reduction of [15] are its ancestors. Our case (b) chooses among block, MDM,
 438 away and toward candidates by realised gain, which subsumes the pairwise-versus-FW choice of
 439 blended pairwise conditional gradients [12]; an empirical comparison with that method is left for
 440 future work.

441 **Synthetic results in brief.** Away and pairwise steps converge at $\epsilon = 10^{-6}$ where the toward and
 442 midpoint steps of [2] do not. Screening reduces 10,000 points to 36–1,054 survivors at $\epsilon = 10^{-5}$.
 443 On seven synthetic and small-real problem classes the optimised solver is fastest to a certified solu-

444 tion on six ($1.6\text{--}2.5\times$ ahead of LIBSVM, LIBLINEAR or NuSVC); the exception, MNIST odd-vs-
445 even, is a dense-support L_2 problem. On the four classes where single MDM also converges, blocks
446 and guard add only $0.9\text{--}1.5\times$ in wall-clock; their clear wins are the two classes (tight tolerance,
447 MNIST) where single MDM does not converge at all.

448 **E Numerical verification of the analysis**

449 Lemma 1 was checked on 4,000 random ensembles including heavily clipped and boundary cases
450 (worst slack at machine precision); Proposition 1 on 500 near-orthogonal ensembles; Lemma 3
451 to 10^{-9} at arbitrary iterates; Lemma 7 on instrumented two-sided runs (no drop iteration added a
452 support index, and every non-drop iteration met the gain bound with M replaced by the computable
453 upper bound $\text{diam } V + \text{diam } W$); and Theorem 2 on instances whose optimal face was computed
454 by an exact QP: at every activating round all zero-weight non-face points were screened, and the
455 final working set equalled the optimal face exactly. Over 5,000 instrumented live block steps none
456 increased the objective. All checks are test files in the released code.